



Introduction to "ivy"

Block storage synthetic workload generator with real time dynamic feedback control

Hitachi Data Systems

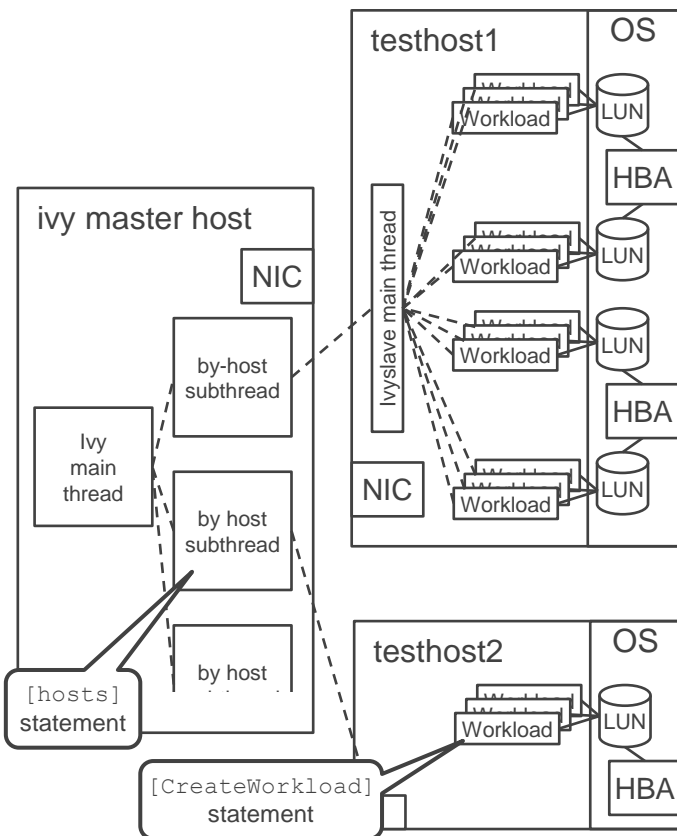
2016-05-16

Allart Ian Vogelesang ian.vogelesang@hds.com +1 408 396 6511

Why ivy?

- Reduce cost
- Speed time to market
- Improve quality
- Simplify test setup
- Measure IOPS at a specified response time (dynamic feedback control)
- Get trust in measurement results through open source

- Written in C++ using Linux native kernel Asynchronous I/O interface
 - Lowest CPU overhead – minimizes test host hardware cost & maximizes measurement accuracy
 - Designed for open source / vendor independence, *with a separately packaged Hitachi proprietary component.*
- Super easy test setup – select by LDEV, by Port, by PG, Pool ID *(by drive type with command device)*
- Can test for minimum time necessary to make a valid measurement to a specified +/- accuracy
- Real-time dynamic feedback control (DFC) enables for the first time to
 - "measure IOPS at 1 ms service time" (using DFC on rolled up host workload data – any vendor's equipment)
 - *"measure IOPS at 50% owning MP core % busy" (using DFC on filtered & rolled up command device data)*
- *Records subsystem command device performance data synchronized with host workload data*
 - *Makes validating / calibrating models easier*
- Records only valid measurement data
 - Checks number of resources reporting (ports, PGs, LDEVs, etc.) and that none running too slowly.
Checks that subsystem does not have any failed components (with command device).
- Enables workflow automation using ivyscript programming language wrapper around ivy engine.



- In ivy, you can layer multiple named “workload” I/O generator threads on LUNs
- Each workload thread uses an I/O sequencer plug-in that generates I/Os in scheduled I/O start time sequence.
 - We try to keep a few I/Os “pre-computed” and ready to go at all times.
 - Generalized interface allows use of any I/O pattern generator, e.g. synthetic or playback.
- Use of C++ with Linux kernel Asynchronous I/O interface means each workload thread can drive any queue depth up to limits of the OS / HBA with state of the art minimal CPU overhead.

ivy – designed for vendor independence

1. Vendor proprietary simple SCSI Inquiry-based LUN discovery tool
 - Vendor's attribute names / values become selectable in ivy.
 - Hitachi LUN discovery tool is open source - https://github.com/Hitachi-Data-Systems/LUN_discovery
2. Ivy
 - Open source - <https://github.com/Hitachi-Data-Systems/ivy>
 - HDS Community site - <https://community.hds.com/groups/ivy>
 - (Send Ian an email to get an invitation to join, but will be public very soon.)
3. For authorized internal Hitachi users with code and license key - vendor proprietary "connector" to retrieve real time storage product configuration and/or performance data
 - Enables ivy dynamic feedback control on internal storage metrics.

- To support any vendor's architecture and terminology in ivy, all you need is a LUN discovery tool that makes a csv file with
 - a header line that defines vendor-specific attribute names, and
 - a data line for each LUN decoding that LUN attributes.
- Hitachi's `LUN_discovery` tool set using "`showluns.sh`" script produces something like this:

```
host,LUN name,HDS product,port,LDEV,PG,CLPR, ...  
testy1,/dev/sdxy,VSP,1A,00:00,1-1,CLPR0, ...  
testy2,/dev/sdyz,VSP,2A,00:01,1-2,CLPR0, ...
```

Functionality expands with proprietary connector

- Vendors are encouraged to prepare their own proprietary "subsystem connector" interface tools to be used with ivy.
 - Collect subsystem configuration data to augment SCSI Inquiry attribute data
 - Select a test configuration based on attributes such as disk drive type.
 - Collect real time subsystem performance data, synchronized and aligned with test host workload data.
 - Facilitates development of modeling tools.
 - Use dynamic feedback control to find the IOPS to reach a target value for
 - subsystem MP % busy, or
 - drive % busy, or
 - cache dirty data % full, etc.

Hitachi's proprietary connector does all those things


- Hitachi's "ivy command device connector" (ivy_cmddev) is not part of the ivy open source project and will remain proprietary for internal Hitachi R&D use.
 - "command device" based functionality is not available to the public.
- Ivy presentation material and sample demo output & videos describe Hitachi command device connector functionality so as to
 - Show what ivy was designed to do
 - Serve to illustrate the functionality that other vendors could achieve by developing their own connectors.

[hosts] statement – configuration discovery

command line: `ivy test2`

`test2.ivyscript`

```
[hosts] "testhost1"  
      to "testhost3"
```



```
Host,LUN,HDS Product,LDEV,PG  
testhost1,/dev/sdxy,VSP,00:00,1-1  
testhost2,/dev/sdxy,VSP,00:01,1-2  
testhost3,/dev/sdxy,VSP,00:02,1-3
```

testhost1

ivyslave
executable

testhost3

testhost2

showluns.sh
executable

Separate vendor-proprietary
SCSI Inquiry tool outputs csv
file decoding storage LUN
attribute names & values

csv column headings
become selectable in ivy

```
Host,LUN,HDS Product,LDEV,PG  
testhost1,/dev/sdxy,VSP,00:00,1-1
```

[CreateRollup] "LDEV" ;

- A rollup is what mathematicians call a "partition" – a way to divide a collection of things (in this case host-LUN-workload threads) into subgroups.
- There is by default always a rollup called "all" which has one rollup instance, also called "all" which consists of all workloads on all LUNs on all hosts.
- Otherwise, a rollup name must be combination of LUN attribute names
- If you have multiple subsystems under test, to get data rolled up by subsystem by port say

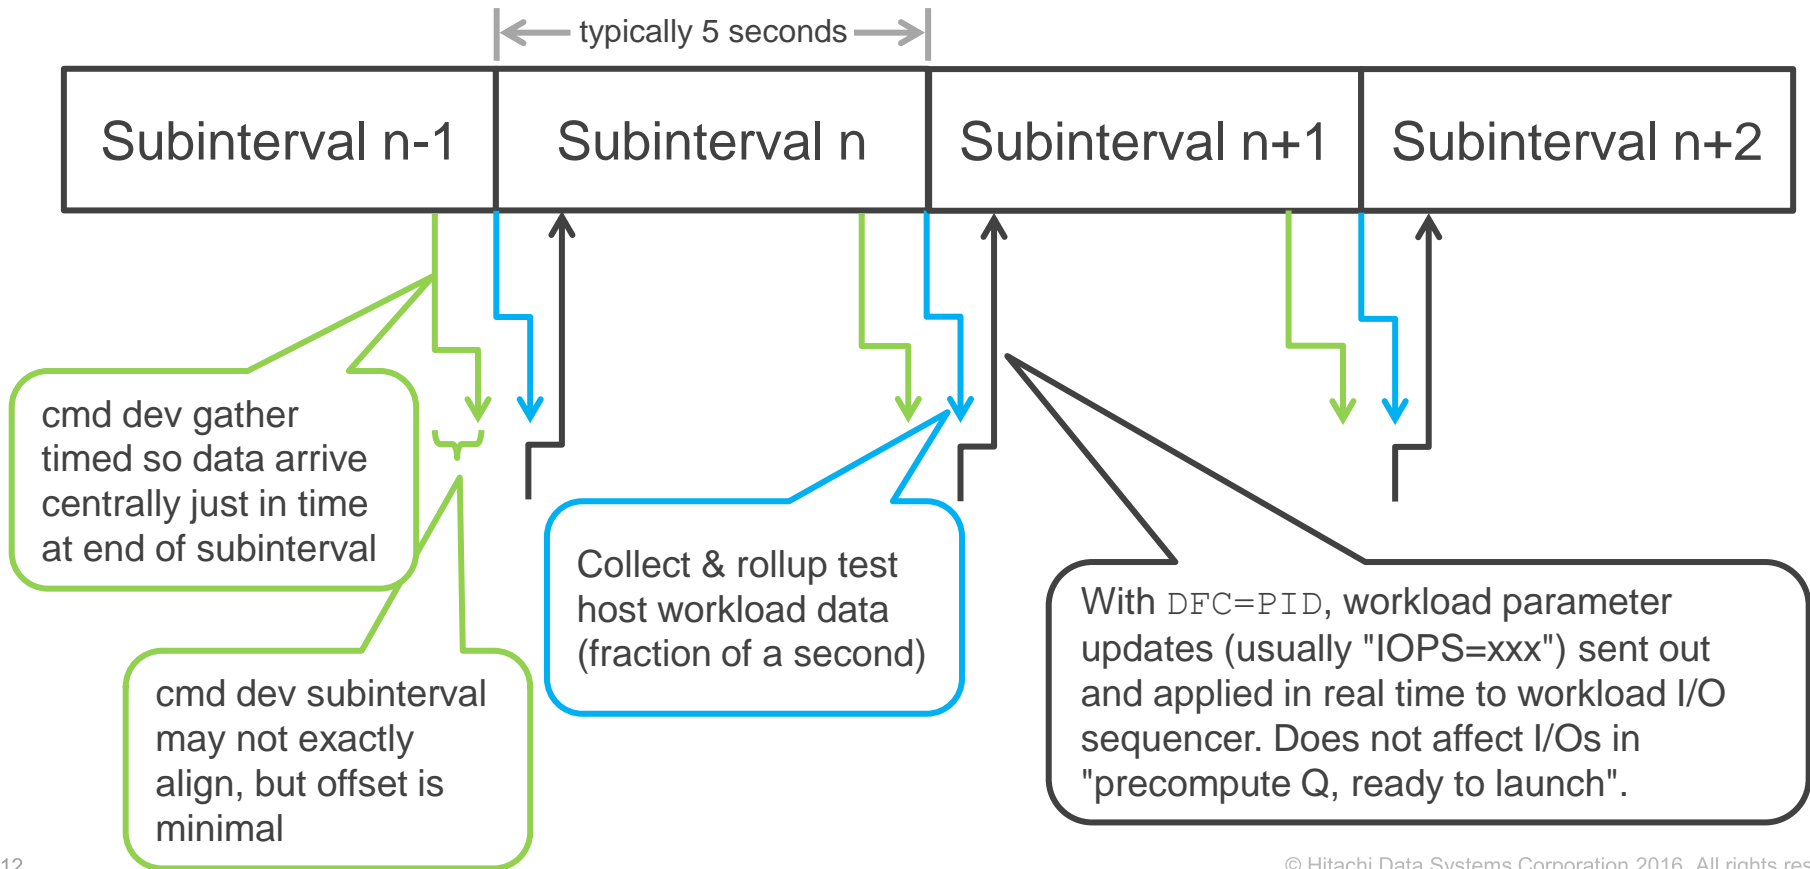
```
[CreateRollup] "serial_number+port";
```

and get instances like "410034+1A".


Rollups are multi-purpose

- The default is to print csv files by rollup instance.
 - Say `[nocsv]` to suppress.
- On the `[Go]` statement to run a subinterval sequence, `focus_rollup="PG"`
 - Both `dfc=pid`, the dynamic feedback control feature, as well as `measure=on`, the "seen enough and stop" `[Go]` statement feature operate at the granularity of the rollup instance.
- `[CreateRollup] "port" [quantity] 32;`
 - Will invalidate measurements if we don't get data for exactly 32 subsystem ports.
- `[CreateRollup] "port" [MaxDroop] "20%";`
 - Invalidates if there is a port whose IOPS is more than 20% slower than the fastest port.


ivy dynamic feedback control (DFC)



industry standard "PID Loop" controller




measure=on, cooldown_by_wp=on




DFC on command device connector data


For parity group busy % = 10%, 20%, 30%, ...



Fine-grained dynamic feedback control




DFC on host workload data – non proprietary




Service time histograms

Read Service Time at 90% of max IOPS




Write Service Time at 90% of max IOPS



Workloads can be placed within the LUN

Sequential MB/s by zone



- May, 2016 initial version.
 - Ready for initial use by Hitachi internally
 - Approved for release as open source release. Now preparing HDS community site for opening for public access by end of May.
 - Demos available for all functionality.
 - No official QA performed (developer testing only), no end user experience yet.
 - Support for VSP Gx00, most, but not all command device data.
 - Support for HUS100 series.

- Need "shakedown" user experience before layering more development
- Dedupe / compression features – in development
- Some ideas for the future:
 - Repackage access to ivy engine via a RESTful API.
 - Build CLI on RESTful API to enable programming ivy in Python or any other language
 - Remove outer ivyscript programming layer wrapper, exposing ivy engine control statements as CLI commands.
 - Develop VSP G1000 support
 - Extend VSP Gx00 support
 - Develop "auto gain control" feature for PID loop
 - Transition to POSIX asynchronous I/O and port to other OS platforms.



Thank You

HITACHI
Inspire the Next 