



# ivy output csv files

## and the ivy output csv file loader Excel tool

**2016-03-30**

Allart Ian Vogelesang [ian.vogelesang@hds.com](mailto:ian.vogelesang@hds.com) +1 408 396 6511

- Prerequisite
  - "introduction to ivy" PowerPoint material.
- Please go over this presentation to learn how to explore the ivy demo output.
- The "programming ivy reference" presentation has full details and can actually be read standalone, but it's recommended to go through this presentation and review the demo video series first.

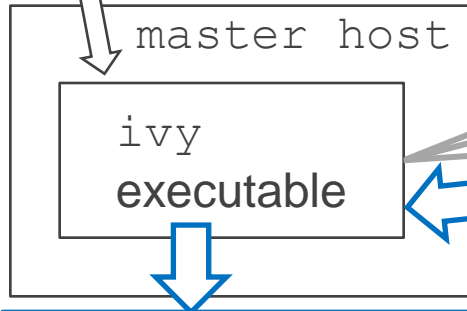
- The `[hosts]` statement specifies the test hosts (I/O driver hosts) to use, and the `[select]` clause in that statement is used to filter "all discovered LUNs" to arrive at "available test LUNs".
- The `[CreateWorkload]` sets up I/O driver threads on test LUNs.
- The `[CreateRollup]` statement sets up two-way structures used to centrally roll up detail data from test host workloads, and to send out I/O driver parameter updates to those workloads.
- The `[Go]` statement runs a subinterval sequence.

# [hosts] statement – configuration discovery

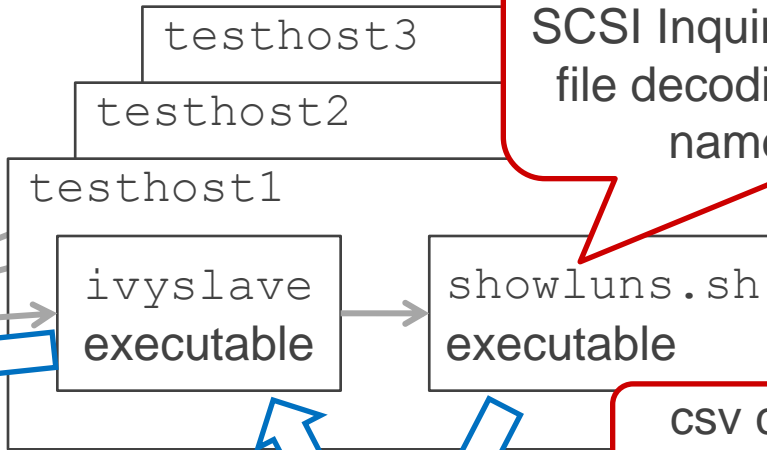
command line: `ivy test2`

`test2.ivyscript`

```
[hosts] "testhost1"  
      to "testhost3"
```



```
Host,LUN,HDS Product,LDEV,PG  
testhost1,/dev/sdxy,VSP,00:00,1-1  
testhost2,/dev/sdxy,VSP,00:01,1-2  
testhost3,/dev/sdxy,VSP,00:02,1-3
```



Separate vendor-proprietary  
SCSI Inquiry tool outputs csv  
file decoding LUN attribute  
names & values

csv column headings  
become selectable in ivy

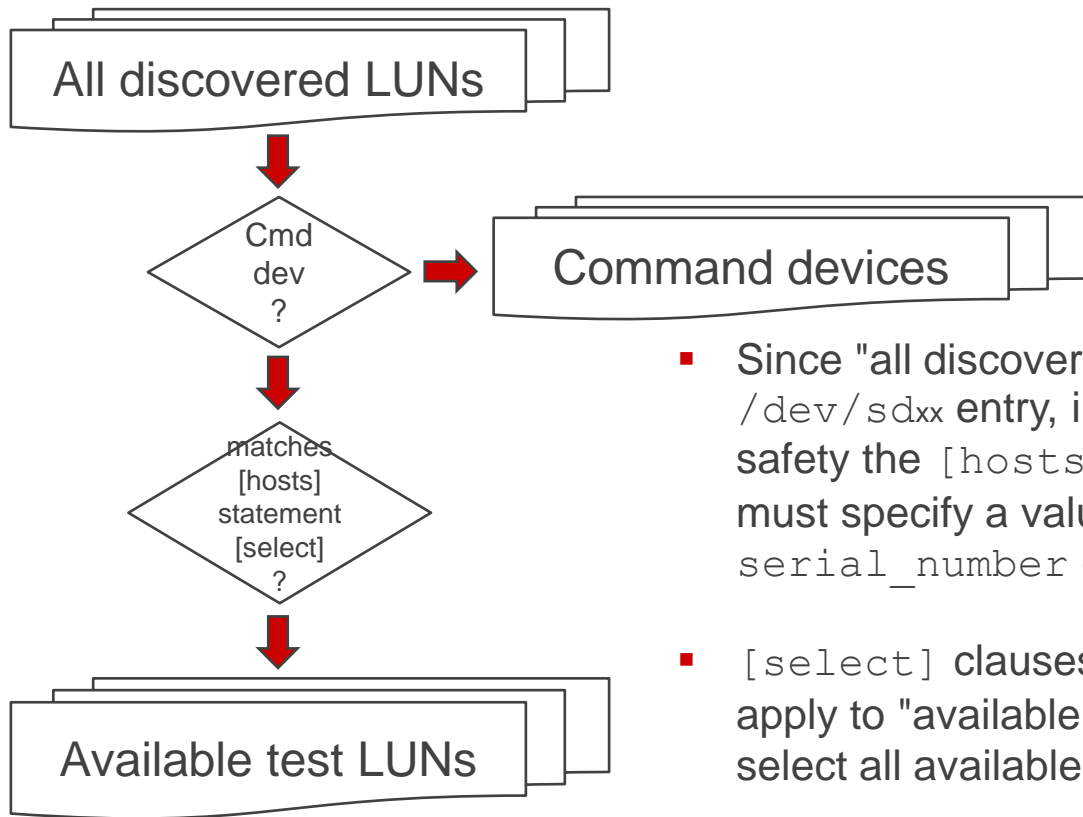
```
Host,LUN,HDS Product,LDEV,PG  
testhost1,/dev/sdxy,VSP,00:00,1-1
```

# "All discovered LUNS"

```
Host,LUN,HDS Product,LDEV,PG  
testhost1,/dev/sdxy,VSP,00:00,1-1  
testhost2,/dev/sdxy,VSP,00:01,1-2  
testhost3,/dev/sdxy,VSP,00:02,1-3
```

- The "showluns.sh" csv files from all the test hosts are combined
- Each data line is loaded into a "LUN" object where for each column in the csv file, we file the data value into the LUN under the attribute name taken from the corresponding header line column.
- Later on, if we find out more information about this LUN, using a command device, we may fill in more attribute values.
- **KEY ivy CONCEPT:** `ivyscript [select]` clauses operate on LUN attributes.

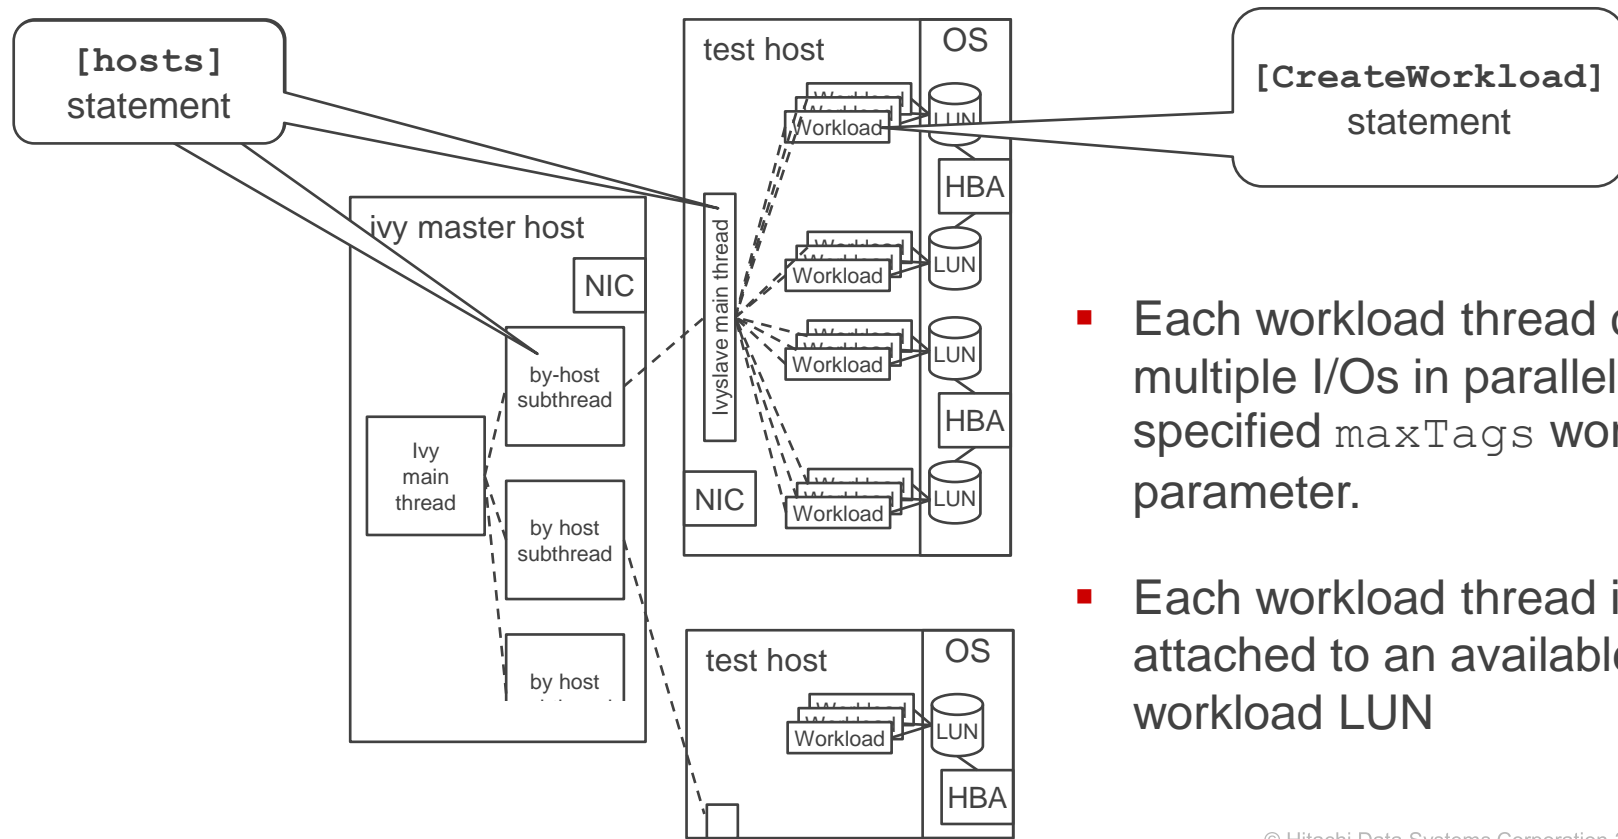
# Available test LUNs



- Since "all discovered LUNs" contains every `/dev/sdxx` entry, including test host boot drives, for safety the `[hosts]` statement `[select]` clause must specify a value for at least one of `serial_number` or `vendor`.
- `[select]` clauses for other ivyscript statements apply to "available test LUNs" and the default is to select all available test LUNs.

- For every Available Test LUN that is mapped to a Hitachi RAID subsystem LDEV for which we have a command device:
  1. We augment the Available Test LUN SCSI Inquiry attributes with the RMLIB API attributes of the underlying LDEV
    - Enables things like `[select] "drive_type = \"NFH1B-P3R2SS\""`
  2. Using the RMLIB API configuration data and using static tables of relationships for the specific subsystem model, we further augment the LUN with "indirect" attributes.
    - Tagging the LUN with the names of the associated subsystem configuration elements of different types enables us to filter real-time RMLIB API performance data by the workload LUNs which comprise a rollup instance.

# Setting up test host workload threads



- Each workload thread can issue multiple I/Os in parallel up to the specified `maxTags` workload parameter.
- Each workload thread is attached to an available workload LUN



# The [go] statement launches a "test step"

- [go] starts the workload threads running a series of "subintervals", typically each 5 seconds long.
- At the end of each subinterval, data are rolled up to the ivy master host.
  - Test host workload data
  - Command device data, if a command device was discovered
- The master host examines the data, optionally sends out real-time I/O generator updates to the workload threads, and then decides whether to continue for another subinterval or not.

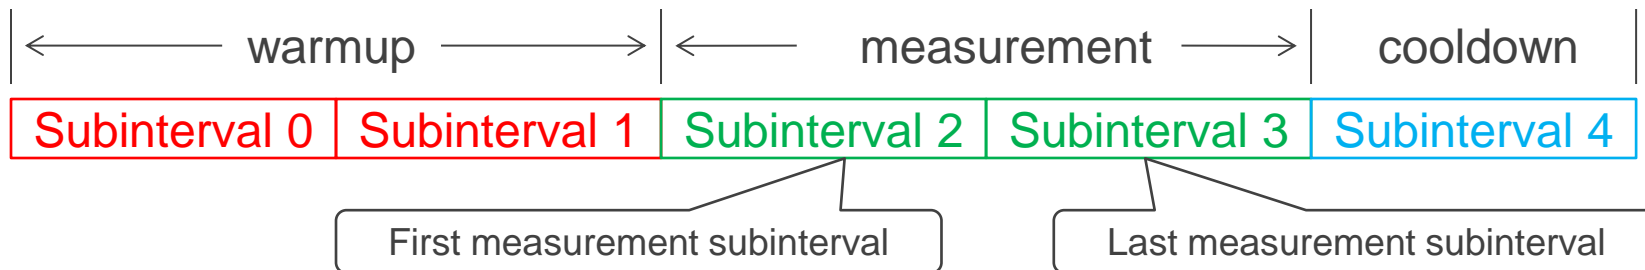
# For each subinterval within a test step ...

- You get a csv line in each csv file within the test step subfolder
- Within the test step subfolder
  - There is a subfolder for each rollup with a csv file for each rollup instance.
    - For example, for a rollup by `host`, you might have csv files for `testhost1`, `testhost2`, and `testhost3`, in the `host` subfolder.
  - There is a subfolder for each subsystem we have a command device for, and within that there is a subfolder for each type of command device resource type for which we collect data.
    - E.g., a csv file for `00:00` in the `LDEV` folder in the `410034` subfolder.

# Test step "success" vs. "failure"

- A test step may run for a fixed number of `warmup_seconds` and `measure_seconds`, in which case the test step is always a "success".
- However, if `measure=on` is specified on the `[Go]` statement, `warmup_seconds` and `measure_seconds` specify minimum times. After that, ivy keeps running more subintervals until it has "seen enough" to make a statistically valid measurement to the required accuracy (success), or until `timeout_seconds` have been reached (failure).
- If the test step is a "success", a summary csv result line is written describing the rollup over the measurement subintervals, otherwise an error message line is written to the summary csv files.
  - Note: Even if a statistically valid measurement was observed, the test step may still be marked a failure if other test configuration validity criteria are not met.

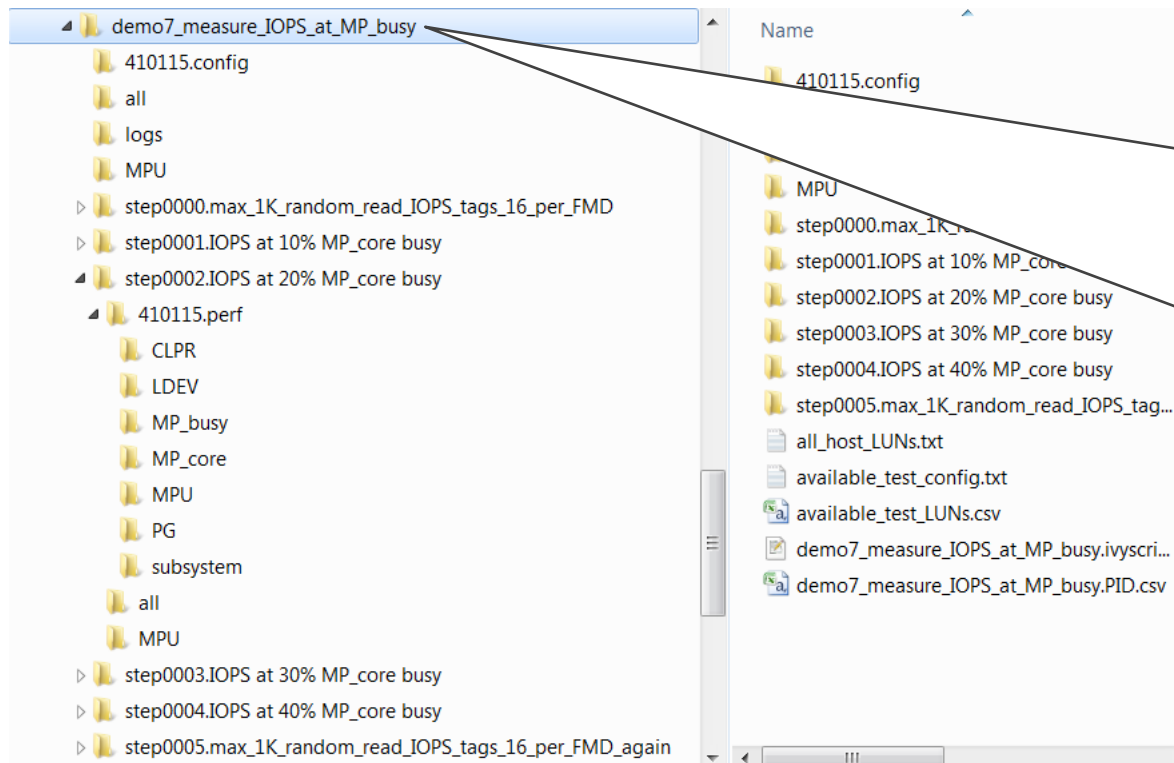
# Success = warmup, measure, cooldown



- For each test step, there is one "measurement" csv line in the "summary" csv files.
- For each rollup, there is a "summary" subfolder in the root test output folder containing a csv file for each rollup instance.
  - The equivalent of the summary.html file in vdbench is the "all=all" csv file in the "all" rollup subfolder.
    - There is always an "all" rollup that has one instance "all" covering all workload threads.
  - With a command device, there are no measurement summary csv files for RMLIB API data like there are for by-subinterval RMLIB API data within a test step subfolder, but some RMLIB API data is filtered by rollup instance, and there are filtered RMLIB API columns within workload csv files.

# Now we are ready to look at ivy output

# ivy output folder structure



The test output folder name is the ivyscript filename with ".ivyscript" removed.

This is also called the "test name", and can be retrieved using the `testName()` ivyscript library function.

This test folder is a subfolder of the `[OutputFolderRoot]` folder, which defaults to `/scripts/ivy/ivyoutput`.

# Files in the test output folder itself

The image shows a file explorer window with two panes. The left pane displays a directory tree for 'demo7\_measure\_IOPS\_at\_MP\_busy'. The right pane shows a list of files in the selected folder. Callouts point to specific files in the right pane:

- All discovered LUNs**: Points to `all_host_LUNs.txt`.
- Particularly with a command device, a concise summary of available test LUNS**: Points to `available_test_LUNs.csv`.
- Available test LUNs**: Points to `available_test_LUNs.csv`.
- A copy of the source .ivyscript program**: Points to `demo7_measure_IOPS_at_MP_busy.ivyscript`.
- A file to help visualize DFC=PID dynamic feedback control**: Points to `demo7_measure_IOPS_at_MP_busy.PID.csv`.

**Directory Tree (Left Pane):**

- demo7\_measure\_IOPS\_at\_MP\_busy
  - 410115.config
  - all
  - logs
  - MPU
  - step0000.max\_1K\_random\_read\_IOPS\_tags\_16\_per\_FMD
  - step0001.IOPS at 10% MP\_core busy
  - step0002.IOPS at 20% MP\_core busy
    - 410115.perf
      - CLPR
      - LDEV
      - MP\_busy
      - MP\_core
      - MPU
      - PG
      - subsystem
      - all
      - MPU
    - step0003.IOPS at 30% MP\_core busy
    - step0004.IOPS at 40% MP\_core busy
    - step0005.max\_1K\_random\_read\_IOPS\_tags\_16\_per\_FMD\_again

**File List (Right Pane):**

Name
410115.config
all
logs
MPU
step0000.max_1K_random_read_IOPS_tags_16_per_FMD
step0001.IOPS at 10% MP_core busy
step0002.IOPS at 20% MP_core busy
step0003.IOPS at 30% MP_core busy
step0004.IOPS at 40% MP_core busy
step0005.max_1K_random_read_IOPS_tags_16_per_FMD_again
all_host_LUNs.txt
available_test_config.txt
available_test_LUNs.csv
demo7_measure_IOPS_at_MP_busy.ivyscript
demo7_measure_IOPS_at_MP_busy.PID.csv

# Subfolders of the ivy test output folder

The screenshot displays a file explorer window with two panes. The left pane shows a tree view of the directory structure, and the right pane shows the contents of the selected folder. Callouts provide explanations for various folders and files.

**Left Pane (Tree View):**

- demo7\_measure\_IOPS\_at\_MP\_busy
  - 410115.config
  - all
  - logs
  - MPU
  - step0000.max\_1K\_random\_read\_IOPS\_tags\_16\_per\_FMD
  - step0001.IOPS at 10% MP\_core busy
  - step0002.IOPS at 20% MP\_core busy
    - 410115.perf
      - CLPR
      - LDEV
      - MP\_busy
      - MP\_core
      - MPU
      - PG
      - subsystem
      - all
      - MPU
    - step0003.IOPS at 30% MP\_core busy
    - step0004.IOPS at 40% MP\_core busy
    - step0005.max\_1K\_random\_read\_IOPS\_tags\_16\_per\_FMD\_again

**Right Pane (Contents):**

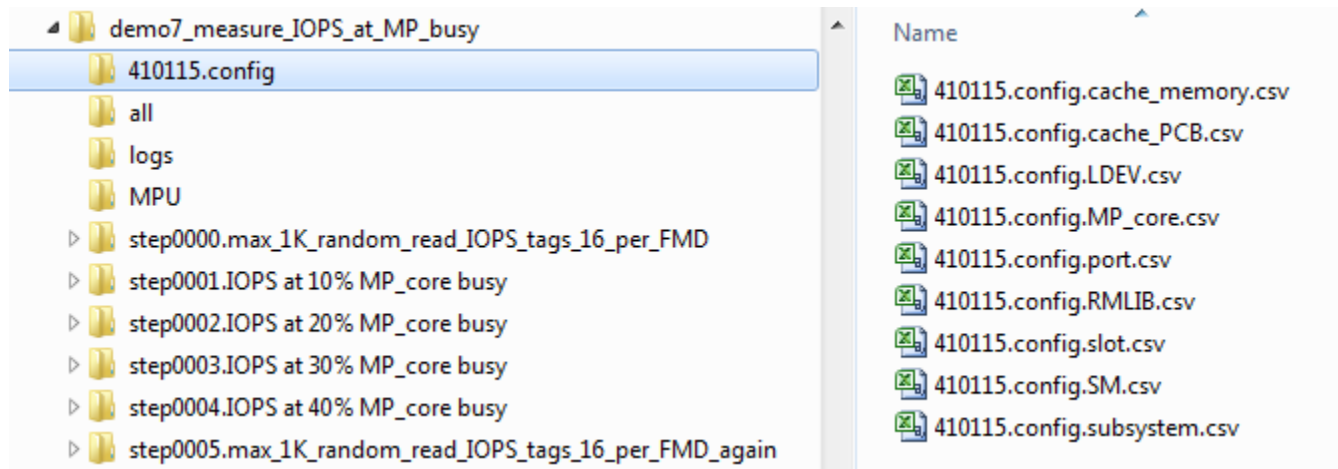
- 410115.config
- all
- logs
- MPU
- step0000.max\_1K\_random\_read\_IOPS\_tag...
- step0001.IOPS at 10% MP\_core busy
- step0002.IOPS at 20% MP\_core busy
- step0003.IOPS at 30% MP\_core busy
- step0004.IOPS at 40% MP\_core busy
- step0005.max\_1K\_random\_read\_IOPS\_tag...
- all\_hosts
- available\_test\_commands.csv
- available\_test\_LUNs.csv
- demo7\_measure\_IOPS\_at\_MP\_busy.ivy
- demo7\_measure\_IOPS\_at\_MP\_busy.perf

**Callouts:**

- If a command device is discovered, a folder containing detailed set of command device configuration csv files is put here
- This is a measurement summary csv folder for the "all" rollup with one summary line for each test step.
- Each ivy process thread on both the master host and the test hosts has its own log file. Use the `-log` command line option to turn on logging of routine events.
- This is a measurement summary csv folder for the "MPU" rollup with one csv file for each MPU containing a summary line for each test step.
- There is a subfolder for each test step containing by-subinterval csv files

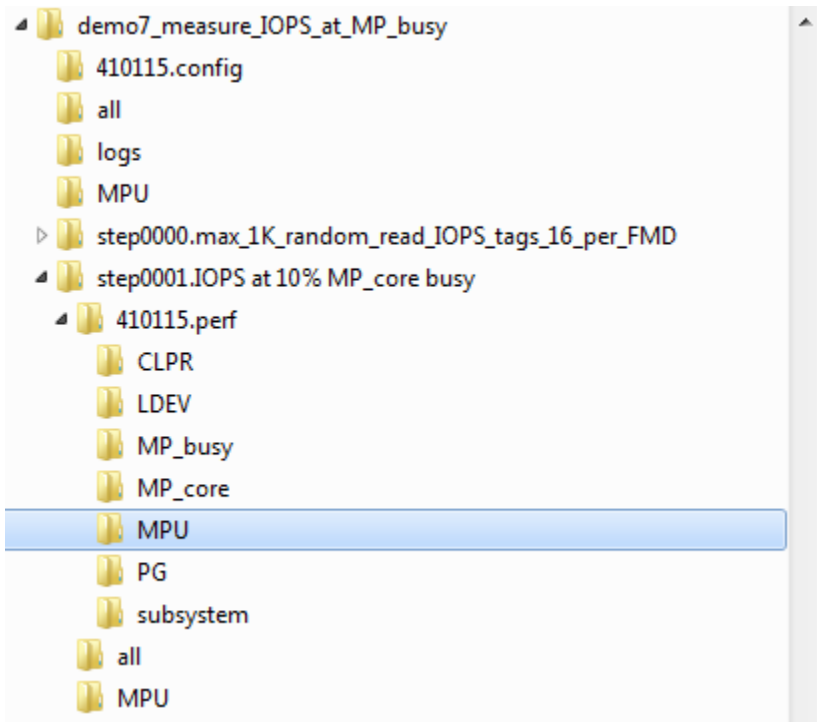


# Ivy records the subsystem configuration



- ... if there is a command device present

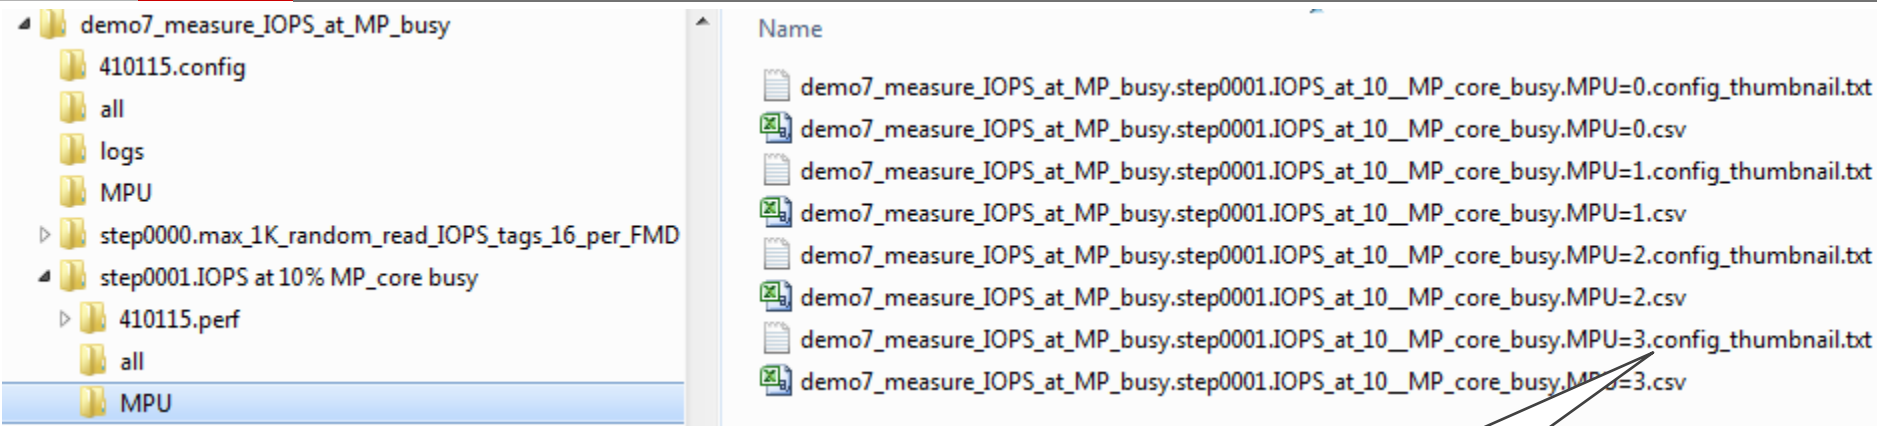
# Ivy records subsystem performance data



Name
demo7_measure_IOPS_at_MP_busy.step0001.IOPS_at_10_MP_core_busy.rmlib.MPU=0.csv
demo7_measure_IOPS_at_MP_busy.step0001.IOPS_at_10_MP_core_busy.rmlib.MPU=1.csv
demo7_measure_IOPS_at_MP_busy.step0001.IOPS_at_10_MP_core_busy.rmlib.MPU=2.csv
demo7_measure_IOPS_at_MP_busy.step0001.IOPS_at_10_MP_core_busy.rmlib.MPU=3.csv

- With a command device, data for the entire subsystem is recorded synchronized with test step subintervals.

# Config thumbnails for each rollup instance



Host-mapped LDEVs by subsystem and drive type or pool ID:

HM800 = 410115, NFH1B-P3R2SS, LDEVs = 00:01 (1)

LDEVs by test host:

172.17.62.15, HM800 = 410115, LDEVs = 00:01 (1)

LDEVs by subsystem port:

HM800 = 410115, port = 3C, LDEVs = 00:01 (1)

Drive counts:

HM800 = 410115, NFH1B-P3R2SS, RAID-5 - 7+1, 8 drives in 1 PG 1-1

HM800 = 410115, NFH1B-P3R2SS total 8 drives.

HM800 = 410115 total 8 drives.

(Drive counts only available  
with a command device)

# There's a tool to automatically load ivy output csv files into an Excel spreadsheet



load\_ivy\_output

# load\_ivy\_output.xlsm

It performs OK ... if you copy your output folder structure to a local SSD drive on your PC and run Excel against that.

**Ivy output csv file loader**

☒ Load ivyscript program TRUE

☐ Load "available testLUNs" csv file FALSE ☒ Load "available\_test\_config.txt" summary TRUE

**RMLIB config csv files**

☐ Load RMLIB config csv files FALSE ☒ LDEV TRUE ☐ MP\_core FALSE ☐ RMLIB FALSE ☐ SM FALSE ☐ Cache FALSE ☒ Port TRUE ☐ Slot FALSE ☒ Subsystem TRUE

☒ Load measurement summary csv files TRUE ☐ Load per instance measurement summary csv files FALSE

☒ Load test overall PID csv file TRUE ☐ Load step PID csv files FALSE

☐ Load measurement summary data validity csv files FALSE

**Detail by subinterval by test step**

☐ Test step number 0 step0000

☒ All test steps 2

**Each test step**

☐ By-subinterval rmlib perf csv files FALSE NOTE: There can be a large number of RMLIB API data csv files, so you may wish to load one type at a time.

☐ CLPR FALSE ☐ LDEV FALSE ☐ MP\_busy FALSE ☐ MP\_core FALSE ☐ MPU FALSE ☐ PG FALSE ☒ subsystem TRUE

☐ by-subinterval rolup csv files FALSE

Select the root folder of the output from your test run.  
For xxxxx.ivyscript, the default output root folder is /scripts/ivy/ivyoutput/xxxxx.

Select an ivy test output root folder

S:\Scripts\ivy\ivyoutput\sample\_output\demoa\_auto\_ranging\_drive\_IOPS\_vs\_response\_time\_DF\_by\_LDEV\

Import

Delete imported data

Makes it easy to look at the data you want to see

Loads both .txt as well as .csv files.

Use this to explore the ivy demo output.



# Questions and Discussion



# Thank You

**HITACHI**  
Inspire the Next 